

Lecture 6: Data Wrangling Continued

Heather Mattie, PhD

September 17, 2026

Recipes of the Day!

Happy Hispanic Heritage Month!

Beef Empanadas



Announcements


- There will **not** be lab this week
- Homework #1 deadline is Friday, September 18
- Grab a duckie!

dplyr::filter() KEEP ROWS THAT satisfy your CONDITIONS

keep rows from... this data... ONLY IF... type is "otter" AND site is "bay"

```
filter(df, type == "otter" & site == "bay")
```

type	food	site
otter	urchin	bay
shark	seal	channel
otter	abalone	bay
otter	crab	wharf



Rmarkdown is magic



Coding question of the day

Using the **gapminder** dataset, calculate the mean fertility rate by region for the year 2015. Which 10 regions had the largest fertility rates that year?

Assume there is no missing data. The following steps may be helpful.

1. Filter the data to only include the year 2015
2. Group by region
3. Calculate the average fertility rate for each region
4. Display the top 10 fertility rates (i.e., the 10 largest fertility rates)

Be sure to run this code first

```
library(dslabs)
library(dplyr)
data(gapminder)
```

Data Wrangling

Data Wrangling Roadmap

1. `dplyr` functions
2. Importing data
3. Reshaping data
4. Combining data frames
5. Parsing dates and times

Data Wrangling

Importing data

Heather Mattie, PhD

Tidy data

Tidy data is a standardized way to structure datasets so they are easy to understand and analyze

- Each observation is a row
- Each variable is a column
- The **tidyverse** package can be used for data wrangling



Gapminder example

Tidy data format

```
library(tidyverse)
library(dslabs)
data("gapminder")
head(gapminder)
```

```
##           country year infant_mortality life_expectancy fertility
## 1      Albania 1960      115.40           62.87           6.19
## 2      Algeria 1960      148.20           47.50           7.65
## 3       Angola 1960      208.00           35.98           7.32
## 4 Antigua and Barbuda 1960      NA           62.97           4.43
## 5      Argentina 1960      59.87           65.39           3.11
## 6      Armenia 1960      NA           66.86           4.55
##  population      gdp continent      region
## 1    1636054      NA    Europe Southern Europe
## 2    11124892 13828152297    Africa Northern Africa
## 3     5270844      NA    Africa  Middle Africa
## 4      54681      NA  Americas      Caribbean
## 5    20619075 108322326649  Americas    South America
## 6     1867396      NA     Asia    Western Asia
```

Gapminder example

Raw data

```
library(tidyverse)
library(dslabs)
filename <- "https://raw.githubusercontent.com/rafalab/dslabs/master/inst/extdata/fertility-two-countries-example.csv"
wide_data <- read_csv(filename)
head(wide_data)
```

```
## # A tibble: 2 × 57
##   country `1960` `1961` `1962` `1963` `1964` `1965` `1966` `1967` `1968` `1969`
##   <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Germany  2.41  2.44  2.47  2.49  2.49  2.48  2.44  2.37  2.28  2.17
## 2 South K... 6.16  5.99  5.79  5.57  5.36  5.16  4.99  4.85  4.73  4.62
## # i 46 more variables: `1970` <dbl>, `1971` <dbl>, `1972` <dbl>, `1973` <dbl>,
## #   `1974` <dbl>, `1975` <dbl>, `1976` <dbl>, `1977` <dbl>, `1978` <dbl>,
## #   `1979` <dbl>, `1980` <dbl>, `1981` <dbl>, `1982` <dbl>, `1983` <dbl>,
## #   `1984` <dbl>, `1985` <dbl>, `1986` <dbl>, `1987` <dbl>, `1988` <dbl>,
## #   `1989` <dbl>, `1990` <dbl>, `1991` <dbl>, `1992` <dbl>, `1993` <dbl>,
## #   `1994` <dbl>, `1995` <dbl>, `1996` <dbl>, `1997` <dbl>, `1998` <dbl>,
## #   `1999` <dbl>, `2000` <dbl>, `2001` <dbl>, `2002` <dbl>, `2003` <dbl>, ...
```

Working directory

The **working directory** is where R will save or look for files by default

- When importing data from a file on a computer, R will search the files in the working directory first

```
getwd()
```

```
## [1] "/Users/hem122/Desktop"
```

Paths

If the file is not in the working directory, a **path** to the file is required to import it

- The path indicates the location (folder) of the file on a computer or on a website

```
path <- system.file("extdata", package = "dslabs")  
path
```

```
## [1] "/Library/Frameworks/R.framework/Versions/4.5-arm64/Resources/  
library/dslabs/extdata"
```

```
filename <- "murders.csv"  
fullpath <- file.path(path, filename)  
fullpath
```

```
## [1] "/Library/Frameworks/R.framework/Versions/4.5-arm64/Resources/  
library/dslabs/extdata/murders.csv"
```

readr

The **readr** package includes functions for reading data stored in a spreadsheet file into R

Function	Format	Typical suffix
read_table	white space separated values	txt
read_csv	comma separated values	csv
read_csv2	semicolon separated values	csv
read_tsv	tab delimited separated values	tsv
read_delim	general text file format, must define delimiter	txt

readxl

The **readxl** package provides functions for reading in Microsoft Excel formats

Function	Format	Typical suffix
read_excel	auto detect the format	xls, xlsx
read_xls	original format	xls
read_xlsx	new format	xlsx

Raw Gapminder data

```
library(tidyverse)
library(dslabs)
filename <- "https://raw.githubusercontent.com/rafalab/dslabs/master/inst/extdata/fertility-two-countries-example.csv"
wide_data <- read_csv(filename)
head(wide_data)
```

```
## # A tibble: 2 × 57
##   country `1960` `1961` `1962` `1963` `1964` `1965` `1966` `1967` `1968` `1969`
##   <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Germany  2.41  2.44  2.47  2.49  2.49  2.48  2.44  2.37  2.28  2.17
## 2 South K... 6.16  5.99  5.79  5.57  5.36  5.16  4.99  4.85  4.73  4.62
## # i 46 more variables: `1970` <dbl>, `1971` <dbl>, `1972` <dbl>, `1973` <dbl>,
## #   `1974` <dbl>, `1975` <dbl>, `1976` <dbl>, `1977` <dbl>, `1978` <dbl>,
## #   `1979` <dbl>, `1980` <dbl>, `1981` <dbl>, `1982` <dbl>, `1983` <dbl>,
## #   `1984` <dbl>, `1985` <dbl>, `1986` <dbl>, `1987` <dbl>, `1988` <dbl>,
## #   `1989` <dbl>, `1990` <dbl>, `1991` <dbl>, `1992` <dbl>, `1993` <dbl>,
## #   `1994` <dbl>, `1995` <dbl>, `1996` <dbl>, `1997` <dbl>, `1998` <dbl>,
## #   `1999` <dbl>, `2000` <dbl>, `2001` <dbl>, `2002` <dbl>, `2003` <dbl>, ...
```

Summary

- Data wrangling is the process of transforming raw data into a useable format called tidy data
- Tidy data format requires each variable to be a column and each observation a row
- The **readr** and **readxl** packages contain functions to read in spreadsheets of data

Data Wrangling

Reshaping data

Heather Mattie, PhD

Raw Gapminder data

```
library(tidyverse)
library(dslabs)
filename <- "https://raw.githubusercontent.com/rafalab/dslabs/master/inst/extdata/fertility-two-countries-example.csv"
wide_data <- read_csv(filename)
head(wide_data)
```

```
## # A tibble: 2 × 57
##   country `1960` `1961` `1962` `1963` `1964` `1965` `1966` `1967` `1968` `1969`
##   <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Germany  2.41  2.44  2.47  2.49  2.49  2.48  2.44  2.37  2.28  2.17
## 2 South K... 6.16  5.99  5.79  5.57  5.36  5.16  4.99  4.85  4.73  4.62
## # i 46 more variables: `1970` <dbl>, `1971` <dbl>, `1972` <dbl>, `1973` <dbl>,
## #   `1974` <dbl>, `1975` <dbl>, `1976` <dbl>, `1977` <dbl>, `1978` <dbl>,
## #   `1979` <dbl>, `1980` <dbl>, `1981` <dbl>, `1982` <dbl>, `1983` <dbl>,
## #   `1984` <dbl>, `1985` <dbl>, `1986` <dbl>, `1987` <dbl>, `1988` <dbl>,
## #   `1989` <dbl>, `1990` <dbl>, `1991` <dbl>, `1992` <dbl>, `1993` <dbl>,
## #   `1994` <dbl>, `1995` <dbl>, `1996` <dbl>, `1997` <dbl>, `1998` <dbl>,
## #   `1999` <dbl>, `2000` <dbl>, `2001` <dbl>, `2002` <dbl>, `2003` <dbl>, ...
```

pivot_longer

The **pivot_longer** function converts wide data into tidy data

- The first argument is which columns to collapse
- The second argument is the name of the column that will store the values of the collapsed columns
- The third argument is the name of the column that will store the values in the original data frame

```
new_tidy_data <- wide_data %>%  
  pivot_longer(`1960`:`2015`, names_to = "year", values_to = "fertility")  
head(new_tidy_data)
```

```
## # A tibble: 6 × 3  
##   country year  fertility  
##   <chr>   <chr>      <dbl>  
## 1 Germany 1960      2.41  
## 2 Germany 1961      2.44  
## 3 Germany 1962      2.47  
## 4 Germany 1963      2.49  
## 5 Germany 1964      2.49  
## 6 Germany 1965      2.48
```

pivot_longer

The `pivot_longer` function converts wide data into tidy data

- The first argument is which columns to collapse
- The second argument is the name of the column that will store the values of the collapsed columns
- The third argument is the name of the column that will store the values in the original data frame

```
new_tidy_data <- wide_data %>%  
  pivot_longer(-country, names_to = "year", values_to = "fertility")  
head(new_tidy_data)
```

```
## # A tibble: 6 × 3  
##   country year  fertility  
##   <chr>   <chr>      <dbl>  
## 1 Germany 1960      2.41  
## 2 Germany 1961      2.44  
## 3 Germany 1962      2.47  
## 4 Germany 1963      2.49  
## 5 Germany 1964      2.49  
## 6 Germany 1965      2.48
```

pivot_longer

The **pivot_longer** function converts wide data into tidy data

- The first argument is which columns to collapse
- The second argument is the name of the column that will store the values of the collapsed columns
- The third argument is the name of the column that will store the values in the original data frame

```
new_tidy_data <- wide_data %>%  
  pivot_longer(-country, names_to = "year", values_to = "fertility") %>%  
  mutate(year = as.integer(year))  
head(new_tidy_data)
```

```
## # A tibble: 6 × 3  
##   country year fertility  
##   <chr>   <int>     <dbl>  
## 1 Germany 1960      2.41  
## 2 Germany 1961      2.44  
## 3 Germany 1962      2.47  
## 4 Germany 1963      2.49  
## 5 Germany 1964      2.49  
## 6 Germany 1965      2.48
```

pivot_wider

The **pivot_wider** function converts tidy data into wide data

- The first argument, **names_from**, indicates which variable will be used as column names
- The second argument, **values_from**, specifies which variable to use to fill in the cells

```
new_wide_data <- new_tidy_data %>%  
  pivot_wider(names_from = year, values_from = fertility)  
  
new_wide_data %>% select(country, `1960`:`1967`)
```

```
## # A tibble: 2 × 9  
##   country    `1960` `1961` `1962` `1963` `1964` `1965` `1966` `1967`  
##   <chr>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 Germany    2.41  2.44  2.47  2.49  2.49  2.48  2.44  2.37  
## 2 South Korea 6.16  5.99  5.79  5.57  5.36  5.16  4.99  4.85
```

Raw Gapminder data #2

```
path      <- system.file("extdata", package = "dslabs")
filename <- file.path(path, "life-expectancy-and-fertility-two-countries-example.csv")

raw_data <- read_csv(filename)
select(raw_data, 1:5)
```

```
## # A tibble: 2 × 5
##   country      `1960_fertility` `1960_life_expectancy` `1961_fertility`
##   <chr>          <dbl>          <dbl>          <dbl>
## 1 Germany         2.41          69.3          2.44
## 2 South Korea      6.16          53.0          5.99
## # i 1 more variable: `1961_life_expectancy` <dbl>
```

Raw Gapminder data #2

```
data <- raw_data %>% pivot_longer(-country)
head(data)
```

```
## # A tibble: 6 × 3
##   country name      value
##   <chr>    <chr>    <dbl>
## 1 Germany 1960_fertility    2.41
## 2 Germany 1960_life_expectancy 69.3
## 3 Germany 1961_fertility    2.44
## 4 Germany 1961_life_expectancy 69.8
## 5 Germany 1962_fertility    2.47
## 6 Germany 1962_life_expectancy 70.0
```

separate

The **separate** function is used to split a single column into multiple columns

- The default character to split on is an underscore **_**

```
data %>% separate(name, c("year", "variable_name"), "_")
```

```
## # A tibble: 224 × 4
##   country year variable_name value
##   <chr>   <chr> <chr>         <dbl>
## 1 Germany 1960 fertility     2.41
## 2 Germany 1960 life         69.3
## 3 Germany 1961 fertility     2.44
## 4 Germany 1961 life         69.8
## 5 Germany 1962 fertility     2.47
## 6 Germany 1962 life         70.0
## 7 Germany 1963 fertility     2.49
## 8 Germany 1963 life         70.1
## 9 Germany 1964 fertility     2.49
## 10 Germany 1964 life         70.7
## # i 214 more rows
```

separate

The **separate** function is used to split a single column into multiple columns

- The default character to split on is an underscore **_**

```
data %>% separate(name, c("year", "variable_name"))
```

```
## # A tibble: 224 × 4
##   country year variable_name value
##   <chr>   <chr> <chr>         <dbl>
## 1 Germany 1960 fertility     2.41
## 2 Germany 1960 life         69.3
## 3 Germany 1961 fertility     2.44
## 4 Germany 1961 life         69.8
## 5 Germany 1962 fertility     2.47
## 6 Germany 1962 life         70.0
## 7 Germany 1963 fertility     2.49
## 8 Germany 1963 life         70.1
## 9 Germany 1964 fertility     2.49
## 10 Germany 1964 life         70.7
## # i 214 more rows
```

separate

The **separate** function is used to split a single column into multiple columns

- The default character to split on is an underscore **_**
- Adding **extra = "merge"** as an argument will ignore future instances of the splitting character

```
data %>% separate(name, c("year", "name"), extra = "merge")
```

```
## # A tibble: 224 × 4
##   country year  name      value
##   <chr>   <chr> <chr>    <dbl>
## 1 Germany 1960  fertility  2.41
## 2 Germany 1960  life_expectancy 69.3
## 3 Germany 1961  fertility  2.44
## 4 Germany 1961  life_expectancy 69.8
## 5 Germany 1962  fertility  2.47
## 6 Germany 1962  life_expectancy 70.0
## 7 Germany 1963  fertility  2.49
## 8 Germany 1963  life_expectancy 70.1
## 9 Germany 1964  fertility  2.49
## 10 Germany 1964  life_expectancy 70.7
## # i 214 more rows
```

separate

The **separate** function is used to split a single column into multiple columns

- The default character to split on is an underscore **_**
- Adding **extra = "merge"** as an argument will ignore future instances of the splitting character

```
data %>%  
  separate(name, c("year", "name"), extra = "merge") %>%  
  pivot_wider()
```

```
## # A tibble: 112 × 4  
##   country year  fertility life_expectancy  
##   <chr>   <chr>      <dbl>         <dbl>  
## 1 Germany 1960      2.41          69.3  
## 2 Germany 1961      2.44          69.8  
## 3 Germany 1962      2.47          70.0  
## 4 Germany 1963      2.49          70.1  
## 5 Germany 1964      2.49          70.7  
## 6 Germany 1965      2.48          70.6  
## 7 Germany 1966      2.44          70.8  
## 8 Germany 1967      2.37          71.0  
## 9 Germany 1968      2.28          70.6  
## 10 Germany 1969      2.17          70.5  
## # i 102 more rows
```

unite

The **unite** function is used to combine two columns into one

```
var_names <- c("year", "first_variable_name", "second_variable_name")
data %>% separate(name, var_names, fill = "right")
```

```
## # A tibble: 224 × 5
##   country year first_variable_name second_variable_name value
##   <chr>   <chr> <chr>                        <chr>                <dbl>
## 1 Germany 1960 fertility                <NA>                2.41
## 2 Germany 1960 life                    expectancy          69.3
## 3 Germany 1961 fertility                <NA>                2.44
## 4 Germany 1961 life                    expectancy          69.8
## 5 Germany 1962 fertility                <NA>                2.47
## 6 Germany 1962 life                    expectancy          70.0
## 7 Germany 1963 fertility                <NA>                2.49
## 8 Germany 1963 life                    expectancy          70.1
## 9 Germany 1964 fertility                <NA>                2.49
## 10 Germany 1964 life                    expectancy          70.7
## # i 214 more rows
```

unite

The **unite** function is used to combine two columns into one

- The default character to unite on is an underscore **_**

```
data %>% separate(name, var_names, fill = "right") %>%  
  unite(name, first_variable_name, second_variable_name)
```

```
## # A tibble: 224 × 4  
##   country year  name      value  
##   <chr>   <chr> <chr>    <dbl>  
## 1 Germany 1960  fertility_NA  2.41  
## 2 Germany 1960  life_expectancy 69.3  
## 3 Germany 1961  fertility_NA  2.44  
## 4 Germany 1961  life_expectancy 69.8  
## 5 Germany 1962  fertility_NA  2.47  
## 6 Germany 1962  life_expectancy 70.0  
## 7 Germany 1963  fertility_NA  2.49  
## 8 Germany 1963  life_expectancy 70.1  
## 9 Germany 1964  fertility_NA  2.49  
## 10 Germany 1964  life_expectancy 70.7  
## # i 214 more rows
```

unite

The **unite** function is used to combine two columns into one

- The default character to unite on is an underscore **_**

```
data %>% separate(name, var_names, fill = "right") %>%  
  unite(name, first_variable_name, second_variable_name) %>%  
  pivot_wider()
```

```
## # A tibble: 112 × 4  
##   country year  fertility_NA life_expectancy  
##   <chr>   <chr>         <dbl>         <dbl>  
## 1 Germany 1960          2.41          69.3  
## 2 Germany 1961          2.44          69.8  
## 3 Germany 1962          2.47          70.0  
## 4 Germany 1963          2.49          70.1  
## 5 Germany 1964          2.49          70.7  
## 6 Germany 1965          2.48          70.6  
## 7 Germany 1966          2.44          70.8  
## 8 Germany 1967          2.37          71.0  
## 9 Germany 1968          2.28          70.6  
## 10 Germany 1969          2.17          70.5  
## # i 102 more rows
```

unite

The **unite** function is used to combine two columns into one

- The default character to unite on is an underscore **_**

```
data %>% separate(name, var_names, fill = "right") %>%  
  unite(name, first_variable_name, second_variable_name) %>%  
  pivot_wider() %>%  
  rename(fertility = fertility_NA)
```

```
## # A tibble: 112 × 4  
##   country year  fertility life_expectancy  
##   <chr>   <chr>      <dbl>         <dbl>  
## 1 Germany 1960      2.41          69.3  
## 2 Germany 1961      2.44          69.8  
## 3 Germany 1962      2.47          70.0  
## 4 Germany 1963      2.49          70.1  
## 5 Germany 1964      2.49          70.7  
## 6 Germany 1965      2.48          70.6  
## 7 Germany 1966      2.44          70.8  
## 8 Germany 1967      2.37          71.0  
## 9 Germany 1968      2.28          70.6  
## 10 Germany 1969      2.17          70.5  
## # i 102 more rows
```

Summary

- The `pivot_longer` function converts wide data into tidy data
- The `pivot_wider` function converts tidy data into long data
- The `separate` function splits a single column into multiple columns
- The `unite` function combines two columns into a single column

Data Wrangling

Combining data frames

Heather Mattie, PhD

Multiple data sources

The data needed for analysis may not be in one table, file, or data frame

- Multiple data sources will need to be **combined** into one data frame
- Join functions, binding functions, set operators can be used

Name	City	Country
Lenna	San Francisco	US
Malcom	New York	US
Akiko	Tokyo	Japan

+

Name	City	Country
Lenna	San Francisco	US
Thomas	London	UK
Diane	Chicago	US



Name	City	Country
Malcom	New York	US
Akiko	Tokyo	Japan
Thomas	London	UK
Diane	Chicago	US

Example data frames

```
library(tidyverse)
library(dslabs)
data(murders)

tab1 <- slice(murders, 1:6) %>%
  select(state, population)
tab1
```

##	state	population
## 1	Alabama	4779736
## 2	Alaska	710231
## 3	Arizona	6392017
## 4	Arkansas	2915918
## 5	California	37253956
## 6	Colorado	5029196

```
data(polls_us_election_2016)


tab2 <- slice(results_us_election_2016, c(1:3, 5, 14, 44)) %>%
  select(state, electoral_votes)
tab2
```

##	state	electoral_votes
## 1	California	55
## 2	Texas	38
## 3	Florida	29
## 4	Illinois	20
## 5	Arizona	11
## 6	Alaska	3


Join functions

Join functions can combine two data frames in different ways


- The **first** data frame mentioned in the code is the **left** data frame
- The **second** data frame mentioned in the code is the **right** data frame




`inner_join(x, y)`




`full_join(x, y)`




`left_join(x, y)`



`right_join(x, y)`




`anti_join(x, y)`



`anti_join(y, x)`

left_join



The **left_join** function retains all rows from the “left” data frame and matches rows from the “right” data frame


- Based on a specified common column
- If there is no match, an NA value will be produced

```
left_join(tab1, tab2)
```

```
## Joining with `by = join_by(state)`
```

##	state	population	electoral_votes
## 1	Alabama	4779736	NA
## 2	Alaska	710231	3
## 3	Arizona	6392017	11
## 4	Arkansas	2915918	NA
## 5	California	37253956	55
## 6	Colorado	5029196	NA

left_join




The **left_join** function retains all rows from the “left” data frame and matches rows from the “right” data frame

- Based on a specified common column
- If there is no match, an NA value will be produced

```
left_join(tab1, tab2, by = "state")
```

##	state	population	electoral_votes
## 1	Alabama	4779736	NA
## 2	Alaska	710231	3
## 3	Arizona	6392017	11
## 4	Arkansas	2915918	NA
## 5	California	37253956	55
## 6	Colorado	5029196	NA

left_join



The **left_join** function retains all rows from the “left” data frame and matches rows from the “right” data frame


- Based on a specified common column
- If there is no match, an NA value will be produced

```
tab1 %>% left_join(tab2)
```

```
## Joining with `by = join_by(state)`
```

```
##      state population electoral_votes
## 1  Alabama    4779736             NA
## 2   Alaska     710231              3
## 3  Arizona    6392017             11
## 4 Arkansas    2915918             NA
## 5 California  37253956             55
## 6  Colorado    5029196             NA
```

right_join



The **right_join** function retains all rows from the “right” data frame and matches rows from the “left” data frame


- Based on a specified common column
- If there is no match, an NA value will be produced

```
tab1 %>% right_join(tab2)
```

```
## Joining with `by = join_by(state)`
```

```
##      state population electoral_votes
## 1  Alaska      710231             3
## 2  Arizona     6392017            11
## 3 California  37253956            55
## 4   Texas           NA             38
## 5  Florida           NA             29
## 6 Illinois           NA             20
```

inner_join



The **inner_join** function retains all rows that have information in both data frames


- Based on a specified common column
- Can be thought of as the intersection of the two data frames

```
inner_join(tab1, tab2)
```

```
## Joining with `by = join_by(state)`
```

```
##      state population electoral_votes
## 1  Alaska      710231             3
## 2  Arizona     6392017            11
## 3 California  37253956            55
```

full_join



The **full_join** function retains all rows from both data frames


- Based on a specified common column
- If there is no match, an NA value will be produced

```
full_join(tab1, tab2)
```

```
## Joining with `by = join_by(state)`
```

##	state	population	electoral_votes
## 1	Alabama	4779736	NA
## 2	Alaska	710231	3
## 3	Arizona	6392017	11
## 4	Arkansas	2915918	NA
## 5	California	37253956	55
## 6	Colorado	5029196	NA
## 7	Texas	NA	38
## 8	Florida	NA	29
## 9	Illinois	NA	20

semi_join



The **semi_join** function retains all rows from the “left” data frame for which we have information in the “right” data frame


- Based on a specified common column
- No columns from the “right” data frame are retained

```
semi_join(tab1, tab2)
```

```
## Joining with `by = join_by(state)`
```

```
##      state population
## 1   Alaska    710231
## 2  Arizona    6392017
## 3 California 37253956
```

anti_join



The **anti_join** function retains all rows from the “left” data frame for which we have no information in the “right” data frame

- Based on a specified common column
- No columns from the “right” data frame are retained

```
anti_join(tab1, tab2)
```

```
## Joining with `by = join_by(state)`
```

```
##      state population
## 1  Alabama    4779736
## 2  Arkansas    2915918
## 3  Colorado    5029196
```

Binding

Binding is another way to combine data frames

- Binding functions do not match data frames by a common variable
- If the datasets don't match by the appropriate dimensions, an error occurs

x			y								
A	B	C				E	F	G			
a	t	1				a	t	3			
b	u	2				b	u	2			
c	v	3				d	w	1			

+

A	B	C	E	F	G
a	t	1	a	t	3
b	u	2	b	u	2
c	v	3	d	w	1

=

A	B	C	E	F	G
a	t	1	a	t	3
b	u	2	b	u	2
c	v	3	d	w	1

x			y								
A	B	C				A	B	C			
a	t	1				c	v	3			
b	u	2				d	w	4			

+

DF	A	B	C
x	a	t	1
x	b	u	2
y	c	v	3
y	d	w	4

bind_cols

The **bind_cols** function returns data frames placed side by side as a single table

- More than two data frames can be combined
- All data frames must have the same number of rows
- Does not match by column values

```
library(tidyverse)
library(dslabs)
data(murders)
head(murders)
```

```
##      state abb region population total
## 1  Alabama  AL  South    4779736    135
## 2  Alaska   AK   West     710231     19
## 3  Arizona  AZ   West    6392017    232
## 4  Arkansas AR  South    2915918     93
## 5 California CA  West   37253956   1257
## 6  Colorado CO   West    5029196     65
```

```
tab1 <- murders[, 1:2]
tab2 <- murders[, 3:4]
new_tab <- bind_cols(tab1, tab2)
head(new_tab)
```

```
##      state abb region population
## 1  Alabama  AL  South    4779736
## 2  Alaska   AK   West     710231
## 3  Arizona  AZ   West    6392017
## 4  Arkansas AR  South    2915918
## 5 California CA  West   37253956
## 6  Colorado CO   West    5029196
```

bind_rows

The **bind_rows** function returns data frames placed one on top of another as a single table

- More than two data frames can be combined
- All data frames must have the same number and names of columns
- Does not match by column values

```
tab1 <- murders[1:2,]  
tab2 <- murders[3:4,]  
bind_rows(tab1, tab2)
```

##		state	abb	region	population	total
## 1		Alabama	AL	South	4779736	135
## 2		Alaska	AK	West	710231	19
## 3		Arizona	AZ	West	6392017	232
## 4		Arkansas	AR	South	2915918	93

Set operators


Set operators are another set of commands that are useful for combining data frames

- Here we will see the **intersect**, **union**, **setdiff**, and **setequal** functions in action
- When the **tidyverse** or **dplyr** packages are loaded these operators can be used on data frames rather than just vectors

A	B	C
c	v	3

intersect(x, y, ...)

Rows that appear in both x and y.



A	B	C
a	t	1
b	u	2

setdiff(x, y, ...)


Rows that appear in x but not y.



A	B	C
a	t	1
b	u	2
c	v	3
d	w	4

union(x, y, ...)

Rows that appear in x or y, duplicates removed). **union_all()** retains duplicates.



Use **setequal()** to test whether two data sets contain the exact same rows (in any order).

intersect

The **intersect** function returns the elements that are common to two vectors, or the rows that are the same in two data frames with the same column names

```
intersect(1:10, 6:15)
```

```
## [1] 6 7 8 9 10
```

```
intersect(c("a","b","c"), c("b","c","d"))
```

```
## [1] "b" "c"
```

```
tab1 <- murders[1:5,]  
tab2 <- murders[3:7,]  
intersect(tab1, tab2)
```

```
##           state abb region population total  
## 1    Arizona  AZ   West    6392017    232  
## 2   Arkansas  AR   South   2915918     93  
## 3 California CA   West   37253956   1257
```

union

The **union** function returns all elements that are in two vectors, or all rows that are in two data frames with the same column names

```
union(1:10, 6:15)

## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15

union(c("a","b","c"), c("b","c","d"))

## [1] "a" "b" "c" "d"

tab1 <- murders[1:5,]
tab2 <- murders[3:7,]
union(tab1, tab2)
```

##	state	abb	region	population	total
## 1	Alabama	AL	South	4779736	135
## 2	Alaska	AK	West	710231	19
## 3	Arizona	AZ	West	6392017	232
## 4	Arkansas	AR	South	2915918	93
## 5	California	CA	West	37253956	1257
## 6	Colorado	CO	West	5029196	65
## 7	Connecticut	CT	Northeast	3574097	97



setdiff

The **setdiff** function returns elements that appear in a first vector but not in a second, or all rows that are in a first data frame but not in a second data frame with the same column names

- **setdiff** is not symmetric

```
setdiff(1:10, 6:15)
```

```
## [1] 1 2 3 4 5
```

```
setdiff(6:15, 1:10)
```

```
## [1] 11 12 13 14 15
```

```
tab1 <- murders[1:5,]  
tab2 <- murders[3:7,]  
setdiff(tab1, tab2)
```

```
##      state abb region population total  
## 1 Alabama  AL  South    4779736    135  
## 2  Alaska  AK   West     710231     19
```

setequal

The **setequal** function returns TRUE if two vectors are the same, regardless of order, or if two data frames are equal, regardless of order

```
setequal(1:5, 1:6)
```

```
## [1] FALSE
```

```
setequal(1:5, 5:1)
```

```
## [1] TRUE
```

```
setequal(tab1, tab2)
```

```
## [1] FALSE
```

Summary

- It is common to need to combine multiple datasets into one
- Join functions match data frames by a common variable
- Binding functions do not match data frames by a common variable, but do need appropriate dimensions to match
- Set operators can be applied to both vectors and data frames